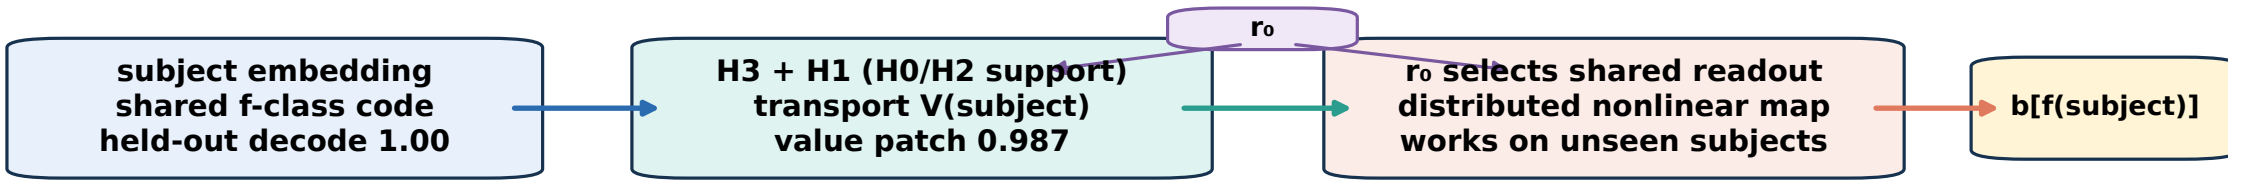
Two learned data structures, two factual-recall algorithms

A Goal 1: shared attribute-class program

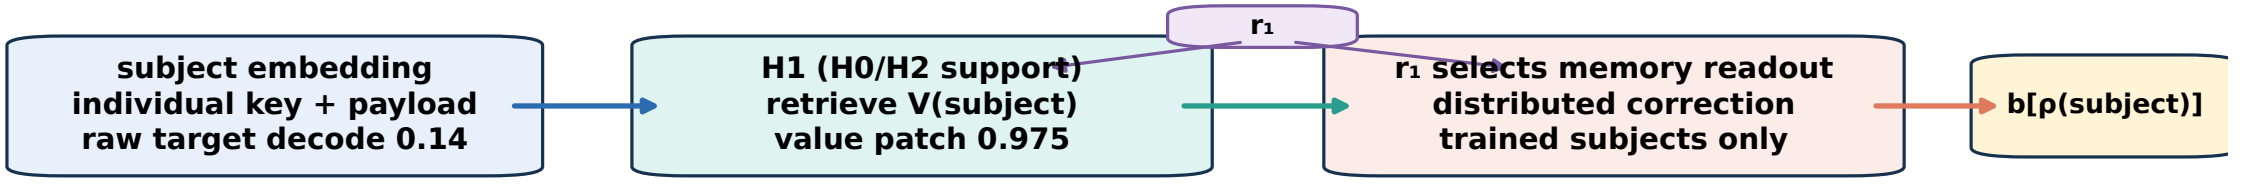
The concept f(subject) is already represented in the embeddings of subjects never seen in relation triples.



Swapping subject values transfers the answer; swapping subject-specific attention weights does not.

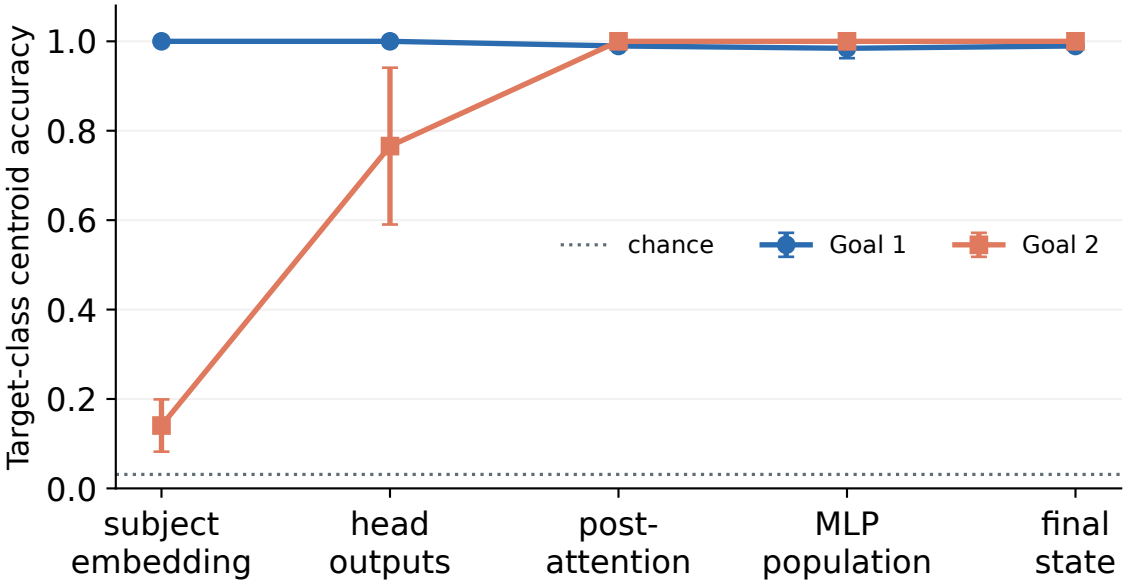
B Goal 2: subject-indexed key–value memory

Repeated triples write an arbitrary target payload into each trained subject record; no shared rule is learned.

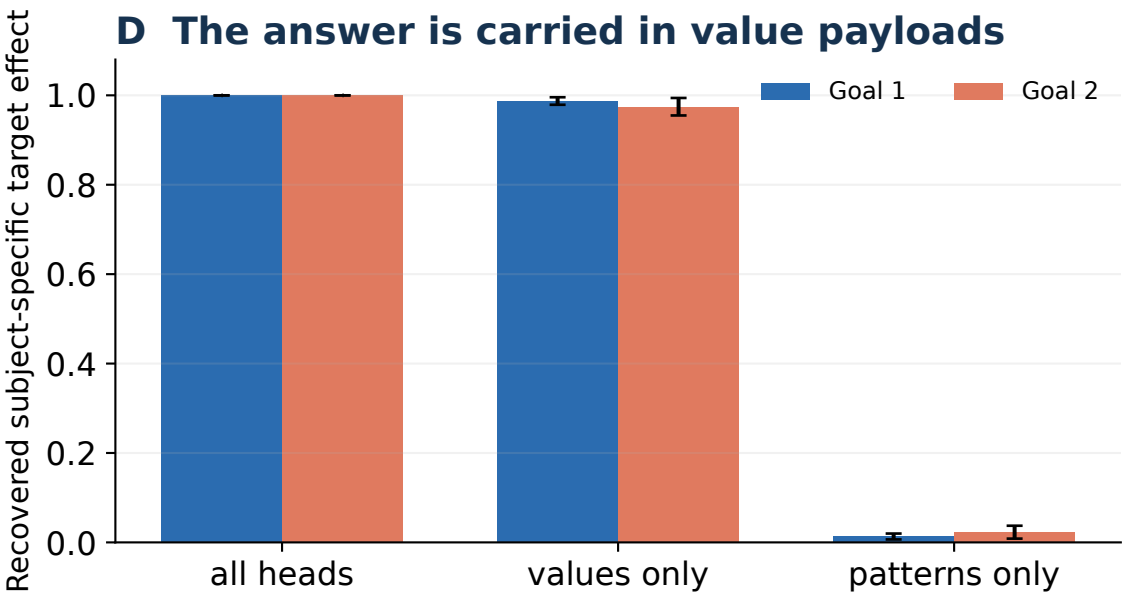


Swapping subject values transfers the answer; swapping subject-specific attention weights does not.

C What is represented, and when?



D The answer is carried in value payloads



Points/bars: mean over seeds 19, 31, 47; error bars: across-seed SD. All decoders and interventions are evaluated on subjects disjoint from discovery.